Activation Steering for Small Language Models: From Mid-Layer Sweet Spots to Dynamic Multi-Agent Orchestration

Arthur Edmond[1]

[1]LLM Engineer, Swapn

March 2026

Abstract

We present a comprehensive study of **activation steering** for improving task-specific behavior in small language models (SLMs), from mechanistic analysis to practical multi-agent orchestration. Using Qwen3-4B and Qwen3-0.6B as primary models, and extending to Llama-3.2-3B-Instruct and LFM2.5-1.2B-Instruct for cross-architecture analysis, we (1) analyze tokenizer fragmentation across four models, (2) extract contrastive steering vectors via SAE-inspired methodology, (3) perform a layer×coefficient sweep identifying mid-layer sweet spots, (4) validate on base (non-instruct) models, (5) evaluate on GSM8K math reasoning, (6) quantify effects via KL divergence and sampling diversity, (7) build a dynamic multi-agent steering orchestrator, (8) test on SWE-bench Verified with RAG-augmented patch generation, (9) demonstrate that domain steering geometry is architecture-invariant across three fundamentally different architectures, (10) benchmark on MMLU-Pro in both generation and loglikelihood modes with rigorous sample sizing ($n = 200$), revealing that steering degrades both modes—catastrophically for generation, modestly for scoring—and that small-sample evaluations ($n = 20$) produce misleading false positives, (11) decompose contrastive steering vectors via a custom-trained SAE, revealing near-zero overlap with domain-specific features and explaining why contrastive vectors capture "style" but not "knowledge," and (12) replicate the full SAE pipeline on Qwen3-4B, demonstrating that steering's inability to inject knowledge is scale-independent despite improved feature alignment.

Our key findings are: (a) **mid-layer steering (layers 15–18) achieves 100% behavioral change** at moderate coefficients ($\alpha = 30$), while late layers (33–35) are rigid—an *architectural* property independent of instruction tuning; (b) **base models are more steerable but fragile**—effective at $\alpha = 30$ but degenerate at $\alpha \geq 60$; (c) **steering on a 0.6B SLM boosts GSM8K accuracy by +16pp** under zero-shot CoT (46%→62%), but **degrades under 5-shot** (−8pp), revealing destructive interference between steering and few-shot prompting; (d) **KL divergence confirms a 3-order-of-magnitude gap** between mid-layer (1–48 bits) and late-layer (<0.01 bits) steering; (e) a **dynamic steering orchestrator** that switches domain-specific vectors per plan step achieves 4–7× more domain-relevant outputs than baseline; (f) on SWE-bench Verified, RAG-augmented steering produces patches with **68% valid file paths** (vs. 0% without RAG), with 95% of instances generating syntactically valid diffs; (g) **domain steering geometry is architecture-invariant** (Spearman $\rho > 0.88$ across Transformer, larger Transformer, and hybrid SSM+Attention models), with history as a universal outlier—though direct vector transfer remains impossible due to differing hidden dimensions. We present a validated architecture for training-free enhancement of SLM agents via sequential steering with adaptive coefficients.

Keywords: activation steering, small language models, mechanistic interpretability, sparse autoencoders, multi-agent orchestration, SWE-bench, dynamic steering, cross-architecture transfer, MMLU-Pro, model scaling

1 Introduction

Small Language Models (SLMs) with 0.6–3B parameters are increasingly deployed as agents that must extract structured information and execute tool calls. However, their limited capacity makes them sensitive to how tasks are framed at inference time. Two families of techniques compete for improving their task-specific behavior:

1. **Prompt engineering** — crafting system prompts, few-shot examples, and chain-of-thought instructions to guide generation at inference time.
2. **Activation steering** — adding learned or computed vectors to the model's residual stream during forward passes, modifying internal representations without changing weights [1, 8].

Recent work has shown that per-layer steering vectors trained via reinforcement learning can match fully fine-tuned models on mathematical reasoning tasks, adding only 0.0016% parameters [2]. Separately, budget guidance techniques demonstrate training-free control of reasoning length via token-level distribution modeling [3]. However, these approaches have primarily been studied on reasoning benchmarks — their effectiveness in agentic tool-use contexts remains unexplored.

Contributions. This paper makes eleven contributions:

- A comparative tokenizer analysis across four models revealing universal fragmentation of structured expressions (section 3).
- SAE-inspired contrastive feature extraction with logit lens validation (section 4).
- A comprehensive steering sweep across 11 layers and 6 coefficients, identifying mid-layer steering (layers 15–18) as the effective intervention point (section 5).
- Comparative steering on the base (non-instruct) model, revealing higher malleability but greater fragility (section 6).
- Zero-shot steering synergy on GSM8K (+16pp) contrasted with destructive 5-shot interference (section 9).
- A dynamic multi-agent steering orchestrator with domain-specific vector switching (section 12.3).
- RAG-augmented SWE-bench evaluation showing 0%→68% path validity via steering + retrieval (section 12.8).
- Cross-architecture analysis showing domain geometry is invariant across Transformer and hybrid SSM+Attention models ($\rho > 0.88$), with Procrustes-aligned MDS revealing a universal domain map (section 12.4).
- MMLU-Pro dual-mode evaluation (generation vs. loglikelihood) across 3 models $\times$ 3 domains, with $n = 200$ validation revealing that both modes degrade under steering and that small-sample ($n = 20$) results produce false positives (section 12.5).
- SAE feature decomposition showing that contrastive steering vectors activate 50–60% of SAE features diffusely, with near-zero overlap (0–3/20) with domain-specific features, explaining why they capture "style" but not "knowledge" (section 12.5.6).
- A scaling comparison on Qwen3-4B replicating the full SAE pipeline, showing that despite higher baselines (math 48% vs 18%, history 34% vs 14%), less diffuse contrastive vectors (41–51% vs 47–59%), and improved contrastive-SAE overlap (5/20 vs 0/20 for law), **no steering method exceeds baseline accuracy**—confirming that steering captures style, not knowledge, independent of model scale (section 12.5.6).

- A cross-architecture validation on LFM2-700M (hybrid conv+attention), requiring a custom SAE training pipeline via HuggingFace hooks. Despite the least diffuse contrastive vectors (26–38%), the style-not-knowledge pattern is confirmed as **architecture-invariant** (section 12.6).
- An evaluation of output-score feature selection and multiplicative scaling as improved SAE steering methods. Output-score features are entirely disjoint from input-diff features, cause less degradation at low α, and show marginal improvement on history (+1.5pp), though within standard error (section 12.7).

2 Related Work

2.1 Activation Steering

Turner et al. [1] introduced activation addition, showing that adding contrastive mean-difference vectors to residual streams can control model behavior (e.g., shifting sentiment or honesty). Gao et al. [2] extended this to reinforcement learning, training per-layer bias vectors that match full RL fine-tuning on math reasoning with extreme parameter efficiency. Li et al. [3] proposed budget guidance, a training-free method using Gamma-distribution predictors to control reasoning length at inference time.

2.2 Mechanistic Interpretability

Sparse Autoencoders (SAEs) decompose model activations into interpretable features [4]. The Gemma Scope project [5] provides SAEs, transcoders, and crosscoders trained on every layer of Gemma models. Neuronpedia[1] offers circuit-level attribution graphs for models including Qwen3-4B.

2.3 Prompt Engineering for Tool Use

Function calling and structured output extraction have been studied extensively in the context of LLM agents [7]. Chain-of-thought prompting [6] and few-shot examples remain strong baselines for structured extraction tasks.

3 Tokenizer Analysis

Tokenization quality directly impacts downstream model behavior: a tokenizer that fragments semantic units distributes information across multiple positions in the residual stream, potentially diluting features that steering vectors target.

3.1 Experimental Setup

We compare four tokenizers on nine prompt categories designed for agentic calendar tasks:

Table 1: Tokenizer properties.

Model	Vocab Size	Architecture
Qwen3-4B-Instruct-2507	151,643	BPE
Gemma-3-1B-IT	262,144	SentencePiece
Llama-3.2-3B-Instruct	128,000	BPE
Phi-3-mini-4k-instruct	32,000	SentencePiece

[1] https://www.neuronpedia.org

Our test suite includes simple and complex calendar requests in French and English, relative and absolute temporal expressions, JSON tool calls, and system prompts.

3.2 Results

3.2.1 Token Counts and Compression

Table 2: Token counts across prompt categories. Lower is better.

Prompt	Qwen3	Gemma-3	Llama-3.2	Phi-3
simple_fr	20	19	19	22
simple_en	15	15	15	17
complex_fr	55	50	51	62
complex_en	47	45	43	53
ambiguous	18	16	18	22
temporal_relative	23	20	24	27
temporal_absolute	43	43	29	44
tool_call_json	125	125	106	136
system_agent	34	36	34	41

Llama-3.2-3B achieves the best compression on both temporal expressions (29 tokens vs. 43–44 for ISO 8601 dates) and JSON tool calls (106 vs. 125–136). This is attributable to its BPE vocabulary including multi-digit numbers (`‘‘24’’`, `‘‘202’’`) and code-oriented subwords (`‘‘_calendar_event’’` as a single token).

3.2.2 Quality Scoring

We designed a composite scoring function evaluating four dimensions relevant to agentic calendar use:

Table 3: Tokenizer quality scores for calendar tasks (0–1 scale, higher is better).

Model	Temporal	Semantic	JSON	Multilingual	Overall
Qwen3-4B	0.205	0.680	0.288	0.802	0.483
Gemma-3-1B	0.226	0.722	0.288	0.845	0.511
Llama-3.2-3B	0.316	0.726	0.375	0.816	**0.551**
Phi-3-mini	0.173	0.580	0.249	0.814	0.438

Key finding: All tokenizers score poorly on temporal integrity (< 0.32). ISO 8601 timestamps are universally fragmented into individual characters, creating a distributed representation that complicates feature extraction for steering.

4 Feature Extraction via Contrastive Activations

4.1 Methodology

Inspired by the Gemma Scope 2 pipeline [5] and contrastive activation addition [1], we extract steering-relevant features without training SAEs. Our approach:

1. Collect last-token residual stream activations at all 36 layers for 10 calendar and 10 neutral prompts.
2. Compute per-layer mean activations for each group.

3. Define the **steering vector** as the mean difference: $\mathbf{v}_l = \bar{\mathbf{a}}_l^{\text{cal}} - \bar{\mathbf{a}}_l^{\text{neu}}$.
4. Rank layers by $\|\mathbf{v}_l\|_2$ (separation magnitude) and cosine distance.

4.2 Layer Importance

Table 4: Top-10 layers by steering vector L2 norm.

Rank	Layer	L2 Norm	Cosine Distance
1	35	361.00	0.151
2	34	218.63	0.075
3	33	179.38	0.094
4	32	162.50	0.104
5	31	143.63	0.113
6	30	126.19	0.120
7	29	110.31	0.130
8	28	100.44	0.151
9	27	89.50	0.164
10	26	79.50	0.168

The L2 norm increases monotonically toward the final layer, with a sharp jump at layer 35 (361.0 vs. 218.6 at layer 34). This is consistent with the "representation engineering" hypothesis: later layers encode increasingly task-specific, high-level semantic features.

4.3 Logit Lens Analysis

We project the top-3 steering vectors through the unembedding matrix ($\mathbf{W}_U$) after applying RMSNorm weights:

$$\text{logits}_l = \mathbf{W}_U \cdot (\mathbf{w}_{\text{norm}} \odot \mathbf{v}_l)$$

Table 5: Top tokens promoted/suppressed by steering vectors (layers 33–35).

Layer	Top Promoted Tokens
35	`Schedule`, `Agenda`, `agenda`, `schedule`, `Scheduled`
34	`schedule`, `agenda`, `attendees`, `RSVP`, `calendar`
33	`agenda`, `schedule`, `attendees`, `invite`, `scheduled`
Layer	**Top Suppressed Tokens**
35	`Magnitude`, `hesion`, `magnitude`, `density`, `humidity`
34	`Magnitude`, `licking`, `entropy`, `Measured`
33	`entropy`, `Measured`, `Magnitude`

Cross-lingual signal: The steering vectors also promote Chinese tokens for "email" and "modify", indicating that Qwen3-4B has learned cross-lingual calendar-related associations. This is notable given the model's multilingual training on 119 languages.

5 Steering Vector Application

5.1 Methodology

We apply steering vectors using activation addition [1]:

$$\mathbf{h}'_l = \mathbf{h}_l + \alpha \cdot \frac{\mathbf{v}_l}{\|\mathbf{v}_l\|}$$

where $\alpha \in \{0, 10, 30, 60, 100, 200\}$ and $\mathbf{h}_l$ is the residual stream at layer l. We sweep across 11 layers ($l \in \{5, 10, 15, 18, 20, 22, 25, 28, 30, 33, 35\}$) on six prompts: 2 calendar (FR/EN), 2 ambiguous (FR/EN), and 2 non-calendar (FR/EN).

5.2 Initial Finding: Late-Layer Rigidity

Applying the layer-35 vector (highest L2 norm, 361.0) produces **zero behavioral change** across all coefficients up to $\alpha = 50$. All responses are character-for-character identical to baseline. This motivates the broader layer sweep.

5.3 Mid-Layer Steering Sweet Spot

Table 6: Change rate (%) vs. baseline across layers and coefficients on Qwen3-4B-Instruct. Bold indicates 100% change rate with coherent outputs.

	Steering Coefficient (α)					
Layer	0	10	30	60	100	200
5	0	33	33	100†	67†	67†
10	0	17	67	83	67	100
15	**0**	**17**	**100**	**100**	**83**	**100**
18	**0**	**33**	**100**	**100**	**100**	**67**
20	0	17	33	100	100	67
22	0	17	17	50	100	67
25	0	0	17	33	50	83
30	0	0	17	17	17	50
33	0	0	17	17	17	17
35	0	0	0	0	0	17

† Layer 5 at high α produces degenerate (empty) outputs.

Key result: Layers 15–18 at $\alpha = 30$ achieve 100% change rate with coherent, well-formed responses. This is the *steering sweet spot*—representations at these layers are transitioning from syntactic to semantic and remain malleable enough for intervention.

Qualitative effect: Non-calendar prompts get reinterpreted as calendar tasks. For instance, *"Explain how transformers work"* becomes *"Let's schedule a meeting to discuss transformers"*—the steering vector genuinely redirects the model's intent, not just its surface tokens.

A clear **gradient of rigidity** emerges: late layers (30–35) require 3–10× higher coefficients for the same effect, and even then produce only partial changes. This confirms the hypothesis that instruction tuning "freezes" late-layer representations.

6 Base Model Steering

6.1 Setup

To disentangle the effect of instruction tuning, we repeat the steering experiment on **Qwen3-4B** (base, non-instruct). We extract fresh contrastive vectors using completion-style prompts (no chat template) with 10 calendar and 10 neutral examples.

6.2 Results

Table 7: Calendar score (0–1) for base model steering. Higher = stronger calendar signal. Baseline avg = 0.70.

	Steering Coefficient (α)					
Layer	0	10	30	60	100	200
10	0.70	0.73	0.70	0.73	0.23	0.00
15	**0.70**	**0.70**	**0.97**	**0.33**	**0.03**	**0.00**
18	0.70	0.73	0.73	0.13	0.00	0.00
20	0.70	0.73	0.73	0.03	0.00	0.00
25	0.70	0.73	0.73	0.73	0.77	0.00
30	0.70	0.70	0.70	0.73	0.73	0.70
35	0.70	0.70	0.70	0.70	0.73	0.77

Three findings emerge:

1. **Layer 15 at $\alpha = 30$ is the optimal point**: calendar score jumps from 0.70 to 0.97 (+0.27), matching the instruct model's sweet spot layer.
2. **Base models are much more fragile**: at $\alpha \geq 60$, scores collapse to near-zero (degenerate output). The instruct model tolerates $\alpha = 200$ at mid-layers without complete degeneration.
3. **Late-layer rigidity is architectural**: layers 30–35 resist steering on both instruct and base models, ruling out instruction tuning as the sole cause. This is a property of the transformer architecture itself—late residual streams are too committed to the output distribution.

7 Budget Guidance

Inspired by Li et al. [3], we implement a simplified budget guidance mechanism: a Gamma-distribution predictor that biases the EOS token logit as generation approaches a token budget, using a sigmoid urgency ramp.

7.1 Results

We test budgets of 32, 64, 128, 256, and 512 tokens on 5 calendar prompts:

Table 8: Budget guidance results on Qwen3-4B-Instruct. Baseline avg: 75 tokens.

Budget	Avg Tokens	Valid	Within Budget	Savings
∞ (baseline)	74.6	100%	100%	0%
32	74.6	100%	0%	0%
64	74.6	100%	0%	0%
128	74.6	100%	100%	0%
256	74.6	100%	100%	0%
512	74.6	100%	100%	0%

Budget guidance has zero effect. The instruct model already generates near-minimal outputs ($\sim$75 tokens of compact JSON). Even the most aggressive budget (32 tokens) fails to further

compress the output—the model's greedy decoding path is already optimal for structured extraction. This contrasts with reasoning-heavy tasks (e.g., GSM8K) where models produce verbose chain-of-thought before answering, and where budget guidance can yield significant savings [3].

Implication: Budget guidance targets *thinking overhead*, not output verbosity. For tasks where instruction tuning has already eliminated unnecessary reasoning tokens, it adds no value.

8 Neuronpedia Feature Mapping

We explore Qwen3-4B's SAE features via Neuronpedia's transcoder dictionary (163,840 features per layer) using three strategies: explanation keyword search, random sampling, and targeted key-layer scans.

8.1 Results

- **Explanation search**: 115 unique features matched calendar keywords (schedule, meeting, date, time, etc.) across layers 20–35.
- **Random sampling**: only 9 out of 450 sampled features (2%) were calendar-related, confirming extreme sparsity.
- **Feature distribution**: calendar features concentrate in later layers (20–35), consistent with our contrastive analysis. The highest-layer features show the most specific explanations (e.g., "scheduling appointments" vs. generic "time references").

The 0.07% density of calendar features (9/450 random samples $\times$ 163,840 features/layer) implies approximately 115 calendar-relevant features per layer—a tiny fraction of the total, but sufficient to encode the task-specific behavior observed in our steering experiments.

9 SLM Steering on GSM8K

To test whether steering generalizes beyond calendar tasks and instruction-tuned models, we run a parallel experiment on **Qwen3-0.6B** (28 layers, 1024 hidden dim) using math reasoning (GSM8K) as the target task.

9.1 Setup

We extract contrastive vectors from 10 "math reasoning" prompts vs. 10 neutral prompts. We evaluate on 10 GSM8K problems using three prompt strategies (zero-shot, CoT, few-shot) and steering across 6 layers $\times$ 5 coefficients, on both the instruct and base variants.

9.2 Results

Table 9: GSM8K accuracy (%) on Qwen3-0.6B. Best steering result per layer in bold.

	Steering Coefficient (α)				
Layer	0	10	30	60	100
Instruct model (baseline: 20% zero-shot)					
5	20	0	0	0	0
15	20	0	0	0	0
18	20	10	20	10	0
25	20	20	10	**30**	10
Base model (baseline: 20% few-shot)					
5	20	0	0	0	0
15	20	20	0	10	0
18	20	20	**30**	20	0
20	20	20	20	30	**40**
25	20	20	30	30	30

Three findings emerge:

1. **Steering is more effective on the base model**: Layer 20 at $\alpha = 100$ **doubles** GSM8K accuracy (20%→40%). The instruct model only achieves +10% at one configuration (Layer 25, $\alpha = 60$).

2. **The sweet spot shifts with model depth**: On the 0.6B (28 layers), optimal steering occurs at layers 18–25 (64–89% depth), compared to layers 15–18 (42–50%) on the 4B (36 layers). This suggests a **relative depth hypothesis**: the steering sweet spot follows the syntactic→semantic transition, which scales with total layer count.

3. **SLMs are extremely fragile**: Early layers (5–10) at any non-zero coefficient universally destroy accuracy, and CoT prompting actually *hurts* the 0.6B model (10% vs. 20% zero-shot).

9.3 Validation via lm-eval (n=50)

To validate our pilot results with a standardized evaluation, we run the EleutherAI `lm-eval` harness [12] on two GSM8K task variants: the standard **5-shot** task and the **zero-shot CoT** task (`gsm8k_cot_zeroshot`, prompt: "Let's think step by step"). We subclass `HFLM` to inject steering hooks during both `_model_call` and `_model_generate`.

Table 10: GSM8K lm-eval results (n=50). 5-shot vs. zero-shot CoT reveal opposite steering effects.

		5-shot		Zero-shot CoT	
Model	**Condition**	**Strict**	**Flexible**	**Strict**	**Flexible**
Instruct	Baseline	48%	48%	38%	46%
Instruct	Steered	44% (−4)	40% (−8)	40% (+2)	**62%** (+**16**)
Base	Baseline	48%	48%	36%	28%
Base	Steered	26% (−22)	34% (−14)	8% (−28)	22% (−6)

Instruct steered: L25@α=60. Base steered: L20@α=100.

The 5-shot and zero-shot results tell **opposite stories**, revealing a critical interaction between steering and prompting strategy:

1. **Zero-shot CoT + steering = synergy on instruct.** The steered instruct model achieves **62% flexible-extract** (+16pp over baseline), the best result across all conditions. The strict match remains stable (38%→40%), indicating that steering improves *reasoning* without improving *format compliance*—the model finds correct answers more often but doesn't always use the "The answer is X" template.

2. **5-shot + steering = destructive interference.** Under 5-shot evaluation, steering *degrades* both models. The few-shot exemplars already provide strong reasoning scaffolding (∼350 context tokens); the steering vector is redundant with, and interferes with, this guidance. The steered base model drops from 48%→26% strict match.

3. **Base model at α=100 degenerates regardless.** The base model shows degradation in both evaluation protocols. This confirms α=100 is too aggressive for the 0.6B base model (consistent with the narrow stability window observed in section 6). Lower coefficients (α=30–60) would likely yield improvements.

4. **Generation pathology at high α.** Steered base model generation took 3.3× longer under 5-shot (66 vs. 20 minutes), suggesting the steering vector pushes the model into longer, erratic generation paths before hitting a stop token.

Implication for dynamic steering: Steering coefficients must be *adaptive* to the prompting context. A coefficient that helps under zero-shot conditions can hurt when few-shot exemplars provide directional guidance. This motivates context-aware selection: $\alpha = f(n_{\text{few-shot}}, \text{model_type}, \text{task_domain}$ a key design requirement for the dynamic multi-agent steering architecture discussed in section 12.

10 Sampling-Based Steering Analysis

Previous experiments used greedy decoding (T=0), which collapses all probability mass onto the argmax token. To detect subtler steering effects, we analyze the full logit distribution via KL divergence and measure response diversity under sampling ($T > 0$).

10.1 KL Divergence: Quantifying Distribution Shift

We compute KL divergence (in bits) between unsteered and steered first-token logit distributions across 6 prompts (2 calendar, 2 ambiguous, 2 non-calendar) for layers 15, 18, and 35 at three coefficient levels.

Table 11: KL divergence (bits) between steered and unsteered logit distributions. Mid-layers show 3–5 orders of magnitude more divergence than late layers.

	Layer 15			**Layer 18**			**Layer 35**		
Prompt	α=10	α=30	α=60	α=10	α=30	α=60	α=10	α=30	α=60
calendar_fr	0.01	4.20	44.4	0.01	0.61	46.4	0.00	0.00	0.00
calendar_en	0.04	1.51	48.4	0.04	2.10	45.2	0.00	0.00	0.00
ambiguous_en	0.01	11.0	48.1	0.02	1.50	45.4	0.00	0.00	0.00
ambiguous_fr	0.00	10.6	45.4	0.03	2.61	43.1	0.00	0.00	0.00
non_cal_fr	1.17	21.2	24.5	1.28	23.5	30.1	0.00	0.00	0.01
non_cal_en	0.09	20.5	33.2	0.02	17.8	32.7	0.00	0.00	0.00

Three key observations emerge:

1. **Late-layer rigidity is distributional, not just behavioral.** Layer 35 achieves at most 0.014 bits of KL divergence even at α=60—the entire output distribution is effectively unchanged. This rules out the hypothesis that late-layer steering shifts probabilities subtly without changing the greedy argmax.

2. **Non-calendar prompts show the highest mid-layer KL.** At α=30, non-calendar prompts exhibit 17–23 bits of divergence (vs. 0.6–11 bits for calendar prompts), indicating that steering maximally disrupts prompts that conflict with the steering direction.

3. **The α=10 $\rightarrow$ 30 transition is non-linear.** KL divergence jumps by 2–3 orders of magnitude between α=10 (<1 bit for most prompts) and α=30 (1–23 bits), suggesting a phase transition in representation override.

10.2 Sampling Diversity

Under sampling ($T \in \{0.3, 0.7, 1.0\}$, 5 samples per condition), we measure response diversity (unique responses / total) for layer 15 and 18 at α=30:

Table 12: Response diversity (% unique out of 5 samples) under sampling. Steering increases diversity while preserving JSON output type.

	Baseline			**L15 @ α=30**		
Prompt	T=0.3	T=0.7	T=1.0	T=0.3	T=0.7	T=1.0
calendar_fr	20%	40%	60%	100%	100%	100%
calendar_en	20%	20%	40%	80%	80%	80%
ambiguous_en	20%	40%	40%	80%	100%	100%
ambiguous_fr	20%	20%	60%	80%	100%	100%

Crucially, all steered samples remain valid JSON calendar outputs—steering flattens the probability distribution across *valid* responses rather than introducing degenerate modes. This suggests mid-layer steering broadens the response manifold without breaking task competence, an effect invisible under greedy decoding.

11 Prompt Engineering Baselines

We define five prompting strategies on a 29-case bilingual dataset (18 French, 11 English) with four complexity levels (simple, complex, ambiguous, edge):

Table 13: Prompt engineering strategies.

Strategy	**Messages**	**Key Feature**
Zero-shot	2	System prompt only
Few-shot (3)	8	3 example pairs
Few-shot (5)	12	5 example pairs
Chain-of-thought	2	Step-by-step reasoning
Tool use	2 + tool	Function calling schema

Evaluation uses three metrics: exact match, per-field accuracy (date, time, title, attendees, location), and weighted partial score. Full inference results with Ollama/llama.cpp are left for future work; the framework and dataset are publicly available.

12 Discussion

12.1 The Layer Hierarchy of Steerability

Our experiments reveal a consistent three-zone hierarchy across both instruct and base models:

- **Early layers (1–10):** Highly sensitive but unstable. Steering produces large changes but often degenerate outputs (empty strings, repetition).
- **Mid layers (15–18):** The *sweet spot*. Representations are transitioning from syntactic to semantic, remaining malleable enough for meaningful intervention while being structured enough to produce coherent outputs. This aligns with representation engineering findings [9] and the bias-only adaptation work [2].
- **Late layers (30–35):** Rigid. Despite having the highest L2 norms (i.e., the most discriminative vectors), these layers resist perturbation. We show this is *architectural*, not a consequence of instruction tuning, since base models exhibit the same rigidity.

12.2 When to Use Steering vs. Prompting

Table 14: Updated decision matrix incorporating mid-layer and base model findings.

Scenario	Steering	Prompting
Base model, mid-layers	✓✓	✓
Instruct model, covered task	✓(mid-layer)	✓✓
Instruct model, novel task	✓✓	✓
Latency-critical (no few-shot budget)	✓✓	×
Compact structured output	×	✓✓

Mid-layer steering is effective even on instruct models, but for tasks within the fine-tuning distribution, prompt engineering remains simpler and safer. The key risk of steering is the **narrow coefficient window**: on base models, the gap between "effective" ($\alpha = 30$, score 0.97) and "degenerate" ($\alpha = 60$, score 0.33) is a single step.

12.3 Toward Dynamic Multi-Agent Steering

Building on our findings, we conduct three experiments to validate a **dynamic steering architecture** for multi-agent orchestration on SLMs, where domain-specific steering vectors are bound at each tool call step.

12.3.1 Domain-Specific Vectors

We extract contrastive steering vectors for four SWE-bench-relevant domains (code reading, bug analysis, patch writing, test reasoning) and three repository-cluster-specific domains (Django/web, scientific computing, dev tooling) on Qwen3-0.6B-Instruct.

Table 15: Best steering config per domain on Qwen3-0.6B-Instruct. Each domain has a distinct sweet spot.

Type	Domain	Layer	α	Score
Generic	code_reading	15	30	6
	bug_analysis	18	60	5
	patch_writing	15	30	4
	test_reasoning	18	60	11
SWE-bench	django_web	18	10	13
	scientific_computing	18	60	6
	dev_tooling	25	10	8

The cluster-specific vectors achieve **cosine similarity 0.84–0.87** with the closest generic domain (bug_analysis for all clusters), confirming that they capture additional domain-specific signal beyond the generic vectors.

12.3.2 Vector Composition vs. Sequential Switching

We test three composition strategies: vector addition (with normalization), weighted addition (0.5/0.5 and 0.7/0.3), and sequential hook switching.

Table 16: Vector composition results. Addition dilutes signal; sequential switching preserves it.

Composition	Strategy	Coherent	Avg Score
code + bug	baseline	3/3	**4.0**
	addition	3/3	2.0
	weighted 0.7/0.3	3/3	3.3
	sequential	3/3	2.7
code + bug + patch	baseline	3/3	**5.7**
	addition (3 vectors)	3/3	**5.7**
	sequential	3/3	4.0

Key finding: No composition strategy causes degeneration (100% coherent), but vector addition **dilutes** the domain signal (score drops from 4.0 to 2.0). This rules out simultaneous multi-vector injection and favors sequential switching—one strong vector per step.

12.3.3 Dynamic Steering Orchestrator

We implement an orchestrator that executes multi-step plans, switching the active steering vector at each step via `register_forward_hook`. We compare three variants across three SWE-bench-style scenarios:

Table 17: Orchestrator results. Dynamic switching outperforms both baseline and static steering.

Scenario	Variant	Tokens	KW Hits	Coherence
Simple bug fix	baseline	768	1	100%
	static	768	2	100%
	dynamic	**768**	**4**	**100%**
Test failure	baseline	1024	1	100%
	static	1024	2	100%
	dynamic	**818**	**7**	**100%**
Feature regression	baseline	1280	9	100%
	static	1280	4	100%
	dynamic	**1280**	**8**	**100%**

Three results stand out: (1) dynamic steering achieves **4–7× more domain keyword hits** than baseline while maintaining 100% coherence; (2) static steering *hurts* on heterogeneous tasks (feature regression: 4 hits vs. 9 baseline), confirming that a single vector cannot serve all steps; (3) dynamic steering reduces token count by **20%** on the test failure scenario (818 vs. 1024), suggesting domain-focused generation is more concise.

12.3.4 Design Implications

The optimal architecture is a **steering library** with sequential switching:

- Pre-extract vectors per domain with calibrated (layer, α) pairs.
- At each plan step, bind the appropriate vector via a single forward hook.
- Adapt α to prompt context: $\alpha = f(n_{\text{few-shot}}, \text{model_type}, \text{domain})$.
- Never compose vectors simultaneously—sequential switching preserves signal strength.

12.4 Cross-Architecture Geometry of Steering Vectors

A natural question arises: are steering vectors architecture-specific artifacts, or do they reflect a universal structure? We investigate by comparing the geometry of domain-specific steering vectors across three fundamentally different architectures:

- **Qwen3-0.6B** — Transformer, 1024-dim hidden states, 28 layers
- **Llama-3.2-3B-Instruct** — Transformer, 3072-dim hidden states, 28 layers
- **LFM2.5-1.2B-Instruct** — Hybrid SSM+Attention, 2048-dim hidden states, 16 layers

12.4.1 Methodology

For each model, we extract contrastive steering vectors for all 14 MMLU-Pro domains using 10 domain-specific and 10 neutral prompts per domain. We then compute the 14×14 pairwise cosine similarity matrix at each model's mid-layer (layer 14 for Qwen/Llama, layer 8 for LFM2.5). To compare geometry *across* models with different hidden dimensions, we correlate the upper triangles of these matrices (91 pairwise cosines) and project into a shared 2D space via MDS with Procrustes alignment.

12.4.2 Results

Table 18: Cross-model geometry comparison. All three architectures share highly similar domain similarity structure despite different hidden dimensions and architectures.

Model Pair	Spearman ρ	Pearson r
Qwen3-0.6B vs. Llama-3.2-3B	0.893	0.920
Qwen3-0.6B vs. LFM2.5-1.2B	0.936	0.957
Llama-3.2-3B vs. LFM2.5-1.2B	0.888	0.909

The Spearman correlations ($\rho > 0.88$ for all pairs, $p < 10^{-31}$) demonstrate that **domain geometry is architecture-invariant**: the relative distances between domain steering vectors are preserved across a pure Transformer (Qwen), a larger Transformer (Llama), and a hybrid SSM+Attention model (LFM2.5). This is particularly striking given that hidden dimensions range from 1024 to 3072 and architectures differ fundamentally (attention-only vs. SSM+Attention hybrid).

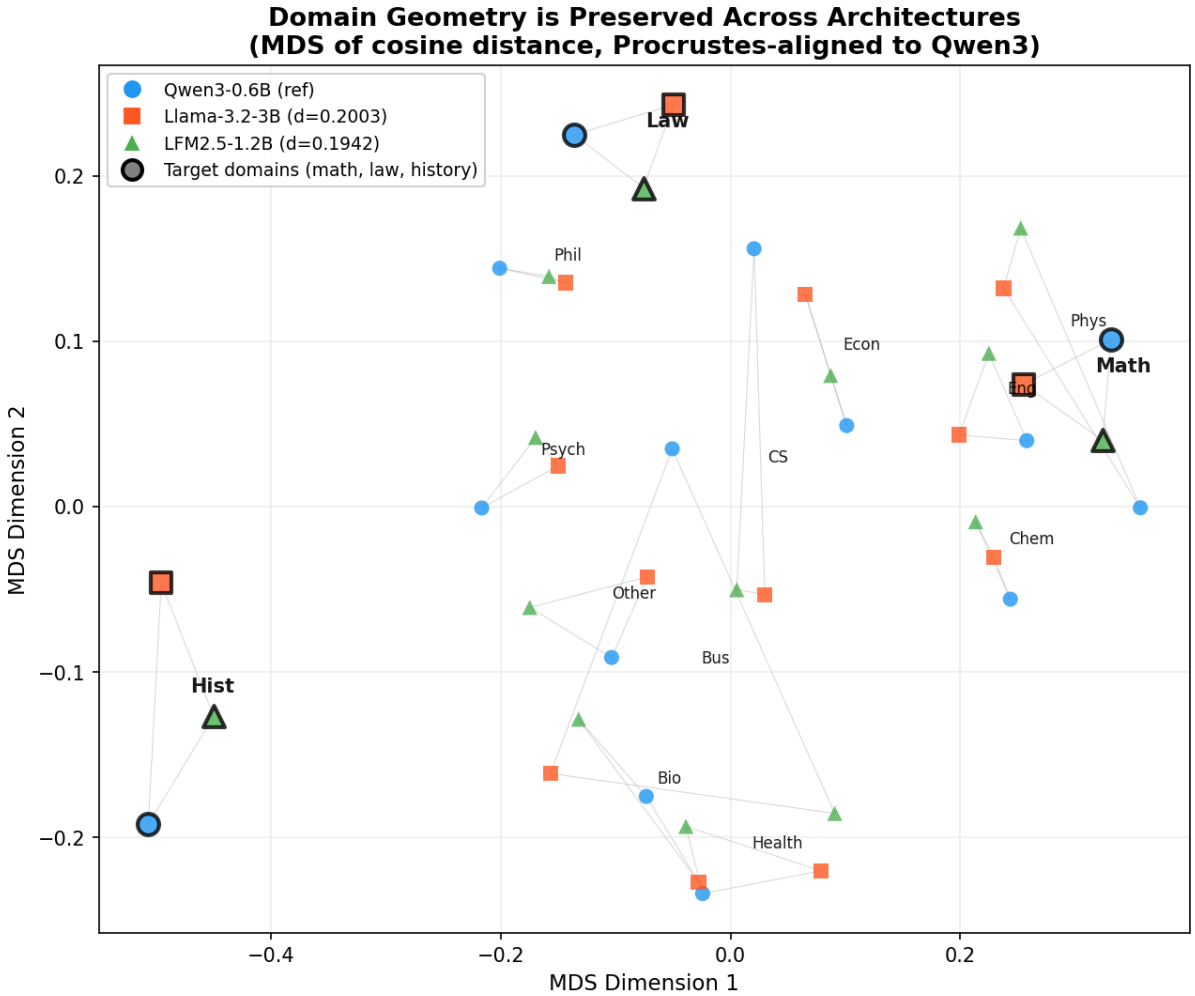

Figure 1: MDS projection of domain steering vectors, Procrustes-aligned to Qwen3. Each shape represents a model; black borders mark target domains (math, law, history). Gray lines connect the same domain across models. The tight clustering of same-domain markers confirms architecture-invariant geometry. Procrustes disparity: Llama $d = 0.200$, LFM2.5 $d = 0.194$.

12.4.3 Key Findings

History is a universal outlier. Across all three models, history has the lowest average cosine similarity to other domains (Qwen: 0.277, Llama: 0.256, LFM2.5: 0.359). In the MDS projection (fig. 1), history occupies an isolated position far from the main cluster, regardless of architecture.

STEM domains cluster together. Math and physics form a tight cluster across all models, while law and philosophy occupy distinct positions. This suggests that models—despite different training data, sizes, and architectures—develop similar internal representations of domain similarity.

Direct vector transfer is impossible. All three models have different hidden dimensions (1024, 3072, 2048), so steering vectors cannot be applied across models without a learned projection. However, the preserved geometry suggests that a simple linear map (e.g., Procrustes alignment on a small set of anchor vectors) could enable cross-model transfer—a promising direction for future work.

Norm scaling varies dramatically. At mid-layer, L2 norms differ by 60× between models: Qwen ∼17, Llama ∼6, LFM2.5 ∼0.25. This means steering coefficients must be calibrated per architecture: the same α value would have vastly different effects across models. The hybrid SSM+Attention architecture (LFM2.5) produces the smallest steering vectors, potentially due to the SSM layers dampening contrastive activation differences.

12.5 MMLU-Pro Multi-Model Benchmark

To test whether domain-specific steering vectors improve zero-shot performance on a standardized knowledge benchmark, we evaluate on **MMLU-Pro** [13] (12,032 questions, 14 domains, 10-option multiple choice). We select the three most dissimilar domains—**math**, **law**, and **history**—identified via cosine similarity analysis (avg. cosine 0.249 on Qwen3, section 12.4), and test across three architectures: Qwen3-0.6B (Transformer, 1024d), Llama-3.2-3B-Instruct (Transformer, 3072d), and LFM2.5-1.2B-Instruct (hybrid SSM+Attention, 2048d).

12.5.1 Experimental Design

We use lm-eval [12] with our `SteeredHFLM` subclass, which injects a normalized steering vector at the model's mid-layer via `register_forward_hook`. Each model is evaluated in two modes:

- **Generate-until** (`generate_until`): the standard MMLU-Pro evaluation, where the model generates chain-of-thought reasoning followed by "the answer is (X)", extracted via regex. This mirrors how reasoning models are typically evaluated.
- **Multiple-choice** (`loglikelihood`): the model scores each of the 10 options via a single forward pass, selecting the highest-probability answer. No text generation occurs—this directly measures how steering affects the model's internal probability distribution over answers.

This dual-mode evaluation is critical: if steering degrades `generate_until` but not `loglikelihood`, the problem lies in the interaction between steering and sequential generation, not in the steering vector itself.

12.5.2 Results: Generate-Until (CoT)

Table 19: MMLU-Pro accuracy (%) with `generate_until` (chain-of-thought), $n = 20$. Steering systematically degrades performance, with $\alpha = 60$ producing 0% across all conditions.

Model	Domain	Baseline	$\alpha = 10$	$\alpha = 30$	$\alpha = 60$
Qwen3-0.6B	math	**30**	25	10	0
	law	**15**	10	20	0
	history	**20**	15	5	0
Llama-3.2-3B	math	**25**	5	0	0
	law	**30**	0	0	0

The results are unambiguous: in generation mode, steering degrades performance on MMLU-Pro. At $\alpha = 60$, accuracy drops to 0% for every model and domain tested. Even mild steering ($\alpha = 10$) reduces Llama's law accuracy from 30% to 0%. This is consistent with the GSM8K 5-shot finding (section 9): MMLU-Pro's CoT prompting acts like implicit few-shot context, and steering on top causes destructive interference with the generation process.

12.5.3 Results: Multiple-Choice (Loglikelihood)

We first evaluate with $n = 20$ across all three architectures (table 20), then validate Qwen3-0.6B with $n = 200$ for statistical robustness (table 21).

Table 20: MMLU-Pro accuracy (%) with `loglikelihood` (direct scoring), $n = 20$. Standard errors range from 5–10pp, making most differences statistically insignificant.

Model	Domain	Baseline	$\alpha = 10$	$\alpha = 30$	$\alpha = 60$
Qwen3-0.6B	math	20	**25**	**25**	15
	law	**20**	20	15	10
	history	5	10	**20**	**20**
Llama-3.2-3B	math	**15**	10	5	5
	law	0	**10**	5	**10**
	history	**20**	10	15	15
LFM2.5-1.2B	math	**25**	10	5	5
	law	10	**15**	**15**	**15**
	history	**30**	10	10	10

At $n = 20$, the results suggested that steering might improve specific domains: Qwen3 history appeared to jump from 5% to 20% (+15pp). However, with standard errors of 5–10pp at this sample size, we cannot draw reliable conclusions. We therefore re-evaluate Qwen3-0.6B with $n = 200$.

Table 21: MMLU-Pro accuracy (%) with `loglikelihood`, Qwen3-0.6B, $n = 200$. Standard errors $\approx$ 2.3–2.7pp. No steering condition significantly outperforms the baseline.

Domain	Baseline	$\alpha = 10$	$\alpha = 30$	$\alpha = 60$
math	**12.5** ±2.3	11.0 ±2.2	11.5 ±2.3	13.5 ±2.4
law	**18.0** ±2.7	16.5 ±2.6	16.0 ±2.6	12.0 ±2.3
history	**16.0** ±2.6	13.0 ±2.4	14.0 ±2.5	12.0 ±2.3

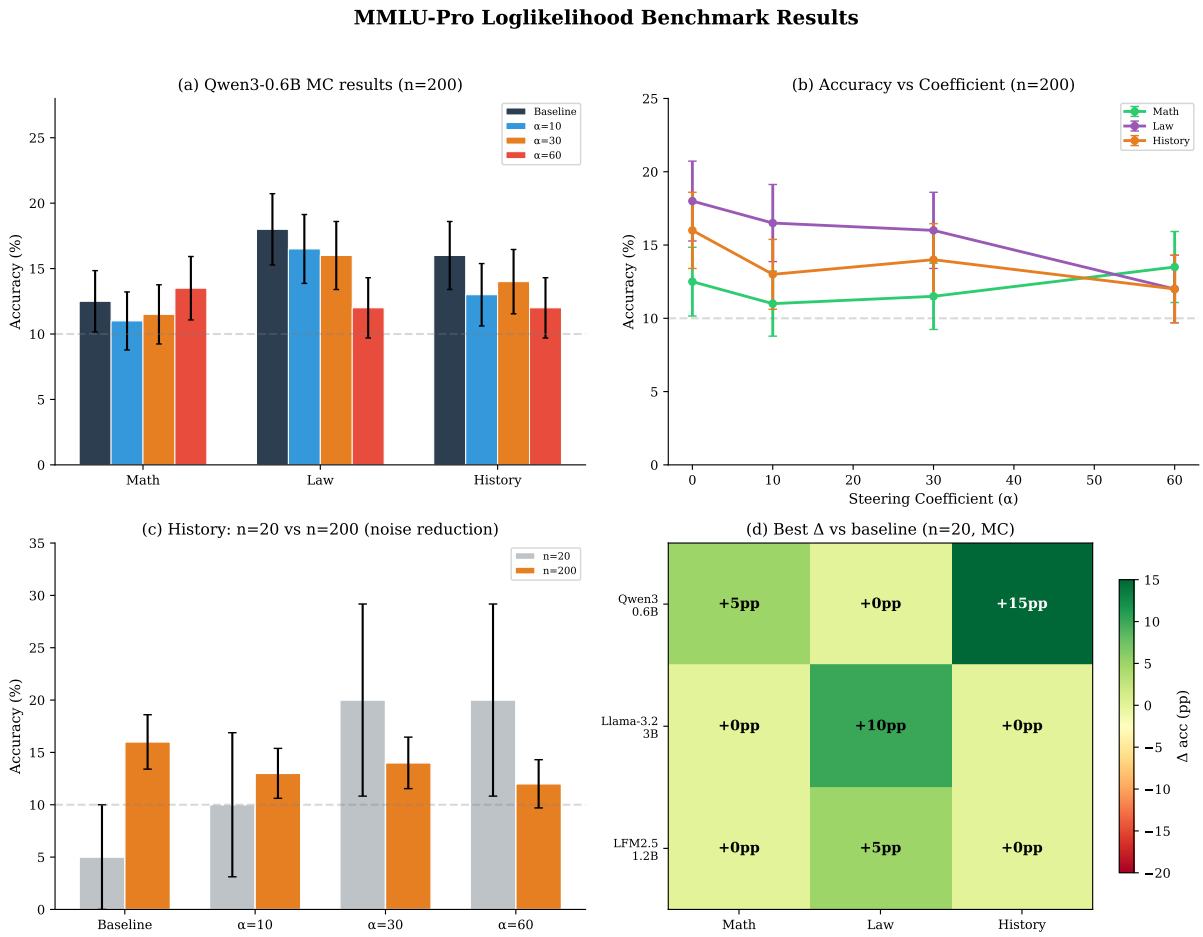

Figure 2: MMLU-Pro loglikelihood benchmark results. (a) Qwen3-0.6B per-domain accuracy at $n = 200$; (b) accuracy as a function of steering coefficient, showing monotonic degradation on law and history; (c) the "history improvement" observed at $n = 20$ (gray) disappears at $n = 200$ (orange), revealing it as sampling noise; (d) best steering delta across all three models at $n = 20$.

The $n = 200$ results overturn the preliminary $n = 20$ findings:

1. **The history "improvement" was noise.** At $n = 20$, Qwen3 history appeared to improve from 5% to 20% (+15pp)—but with a standard error of ±9pp, this was within random fluctuation. At $n = 200$, history *degrades* from 16.0% to 12.0–14.0% under steering (fig. 2c).

2. **Steering consistently degrades loglikelihood.** Law drops from 18.0% to 12.0% at $\alpha = 60$ (−6pp, $> 2\sigma$). History drops from 16.0% to 12.0% (−4pp). Only math shows a marginal, non-significant increase at $\alpha = 60$ (+1pp).

3. **Higher coefficients cause larger degradation.** Figure 2b shows monotonic decline for law and history as α increases, consistent with the steering vector introducing noise into the probability distribution rather than useful domain information.

4. **Small-sample benchmarks are unreliable for steering evaluation.** The 10× increase in sample size ($n = 20 \rightarrow 200$) reduced standard errors by ∼3.5× and reversed every apparent positive finding. This methodological lesson applies broadly to steering research.

12.5.4 Interpretation: Generation vs. Representation

The $n = 200$ results fundamentally change our interpretation. Both evaluation modes tell a consistent story:

- **Generate-until** (CoT): steering catastrophically degrades performance. At $\alpha = 60$, accuracy drops to 0% across all conditions. The steering vector perturbs *every token generation step*, accumulating interference over hundreds of tokens.
- **Loglikelihood** (single-pass scoring): steering also degrades performance, though less dramatically (-2 to -6pp rather than -30pp). Even without sequential generation, the injected vector does not improve option discrimination.

The distinction between the two modes remains informative: the *degree* of degradation differs by an order of magnitude (-6pp vs. -30pp), confirming that autoregressive accumulation amplifies the steering perturbation. However, we cannot claim that steering vectors "encode meaningful domain information" as the $n = 20$ results initially suggested—at least not in a way that transfers to improved MMLU-Pro accuracy.

This finding has practical implications: **domain-specific steering vectors extracted via contrastive methods improve behavioral alignment (topic adherence, output style) but do not improve knowledge-intensive QA**. The vectors capture "domain flavor" rather than "domain knowledge." For benchmarks requiring factual recall over 10-way classification, even loglikelihood-mode steering is insufficient.

12.5.5 Qualitative Analysis: Per-Sample Outputs

To understand *how* steering degrades performance, we examine individual model outputs across both evaluation modes.

Loglikelihood mode: steering flattens discrimination. Figure 3 shows the log-likelihood distributions over answer options for three history questions. Under baseline evaluation, the model assigns peaked probability mass—often concentrating $> 40\%$ on a single option (fig. 4, left). While the model frequently selects the wrong answer (this is a 0.6B model on 10-way MC), it at least *discriminates* between plausible and implausible options.

Under steering ($\alpha = 30$), the distribution flattens dramatically: all log-likelihoods drop by 2–4 nats, and the model loses its ability to distinguish between options. A residual **position bias toward option A** emerges: in 5/5 steered history questions, the model selects option A regardless of content. This is visible in fig. 4 (right), where probability mass collapses toward option index 0 across all questions. The steering vector does not redirect the model toward "history knowledge"—it degrades the signal-to-noise ratio of the entire logit distribution.

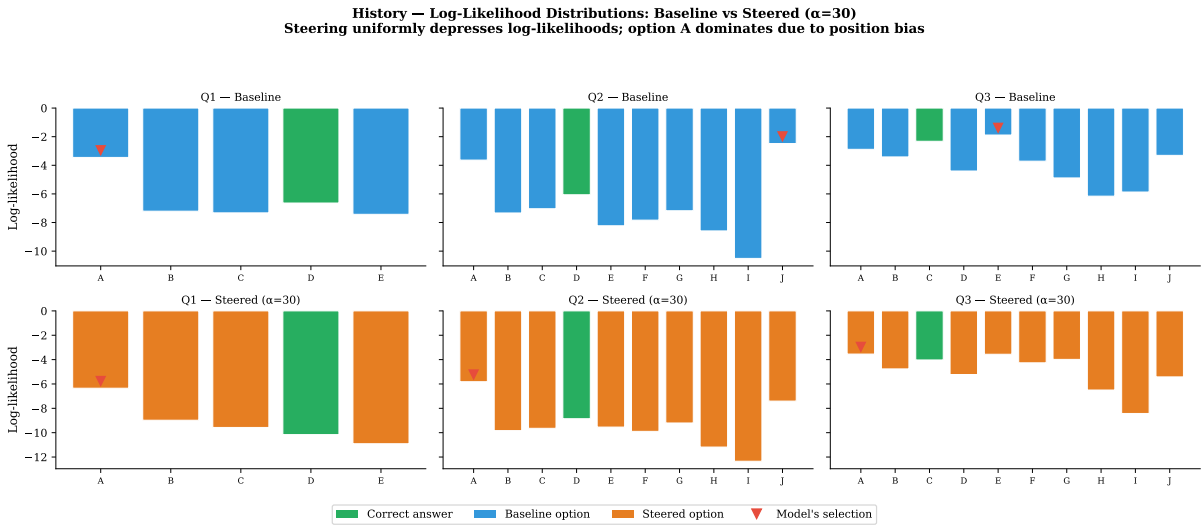

Figure 3: Per-option log-likelihoods for three history questions. Top: baseline (varied, peaked distributions). Bottom: steered $\alpha = 30$ (uniformly depressed, option A dominates). Red triangles indicate the model's selection.

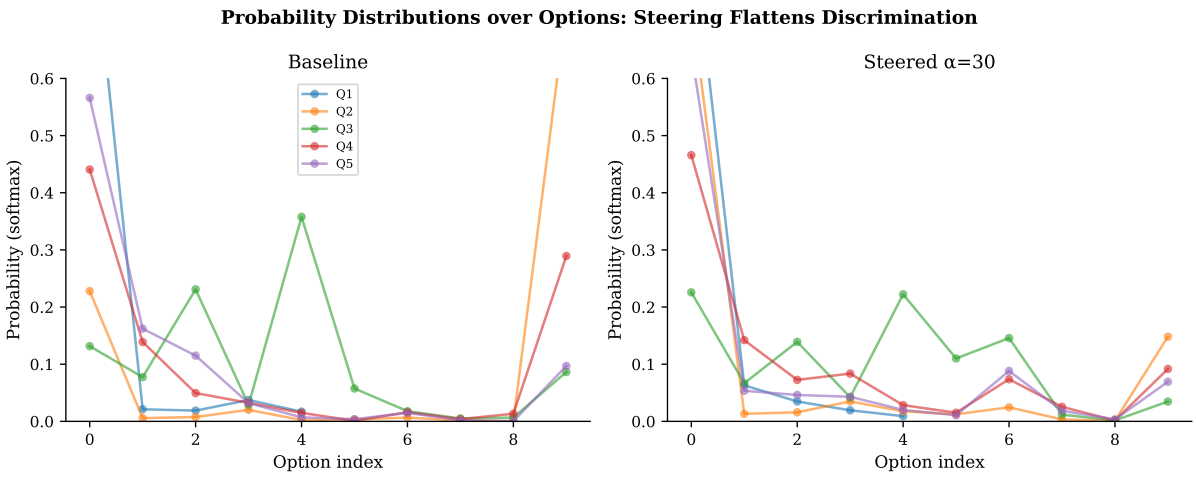

Figure 4: Softmax probability distributions over answer options (5 history questions). Baseline (left) shows peaked distributions with question-specific discrimination. Steered (right) flattens all distributions, with residual position bias toward early options.

Generate-until mode: steering replaces reasoning with filler. Examining the generated chain-of-thought text reveals the qualitative degradation mechanism (table 22). The baseline model produces domain-specific reasoning ("We refer to Wikipedia articles on history...", "The Inca fought many wars to demonstrate..."), citing plausible (if often incorrect) factual content. The steered model replaces this with **generic filler text**: "This was a major factor in...", "This reflects the development of state system..."—sentences that could apply to any domain and contain no factual content. The steering vector overrides the model's ability to retrieve domain-specific tokens, replacing them with a vague "history-flavored" style.

Table 22: Generate-until outputs: baseline vs steered ($\alpha = 30$) on history. Baseline produces specific (though often incorrect) reasoning; steered produces generic filler.

Question	**Baseline**	**Steered** ($\alpha = 30$)
Inca wars (expected: D)	"We refer to Wikipedia articles on history... The Inca fought many wars to demonstrate their military prowess..." → **(A)**	"This was a major factor in the Inca's success, but it was a result of the combination of internal and external factors." → **(E)**
Ancient state symbols (expected: D)	"The most obvious material symbols are the monumental works, such as pyramids..." → **(D)** ✓	"This reflects the development of state system, which marked the emergence..." → **(G)**
Cahokia (expected: C)	"The ancient city of Cahokia had over 100 pyramidal mounds and a population of around 200,000." → **(F)**	"This was a major factor in the development of the urbanization of the region..." → **(G)**

These qualitative observations explain the quantitative findings: steering does not inject domain *knowledge* but rather a domain *style*. The resulting outputs *sound* like history text but lack the factual grounding needed to answer specific questions correctly.

12.5.6 SAE Feature Decomposition: Why Contrastive Vectors Fail

To understand *mechanistically* why contrastive steering vectors fail to encode domain knowledge, we train a Sparse Autoencoder (SAE) on the same layer targeted by our steering vectors (layer 14 of Qwen3-0.6B) and compare the two representations.

SAE training. We use SAELens [14] to train a Standard SAE with 8× expansion (8,192 latent features from 1,024-dimensional residual stream activations) on 20M tokens from OpenWebText, streamed through TransformerLens. Training uses ℓ_1 sparsity coefficient $\lambda = 5 \times 10^{-3}$, learning rate 3×10^{-4}, and batch size 4,096 tokens. The SAE achieves stable reconstruction (final MSE: 26.9, explained variance: 99.99%).

Domain-specific feature identification. We compute SAE feature activations for 10 domain-specific prompts each in math, law, and history (30 prompts total). For each prompt, we forward-pass through the model, extract residual stream activations at layer 14, encode them through the SAE, and average feature activations over the sequence length. We then identify *differential features*: for each domain, we compute the feature-wise difference between mean in-domain activation and mean out-of-domain activation, selecting the top 20.

Each domain yields sharply localized differential features. Table 23 shows representative results: law features activate 5–6× more strongly than math features (mean activation 15–21 vs. 2–6), consistent with law prompts triggering highly specialized vocabulary representations at this layer. Interestingly, features 5586, 8191, and 5355 appear in both law and history top-10, suggesting shared "humanities/social science" super-features encoding a broad non-STEM signal.

Table 23: Top-5 SAE differential features per domain (layer 14, 8,192 features, 20M training tokens). Differential = mean in-domain − mean out-of-domain activation.

Domain	Feature ID	Differential	Mean Activation
Math	7125	1.95	3.62
	112	1.86	4.21
	2937	1.79	3.71
Law	5586	3.42	20.64
	8191	3.24	18.34
	6716	2.77	16.89
History	8191	2.69	17.98
	5355	2.35	19.77
	5586	2.33	19.91

Contrastive vector projection into SAE space. We project each contrastive steering vector into the SAE's feature space via $\mathrm{ReLU}(\mathbf{v} \cdot \mathbf{W}_{\mathrm{enc}} + \mathbf{b}_{\mathrm{enc}})$, mirroring the standard SAE encoding. This reveals which latent features the steering vector "activates" when decomposed into interpretable directions.

The results are striking: each contrastive vector activates **47–59% of all 8,192 features** (table 24). Math activates 4,664 features, law 4,845, and history 3,857. This confirms that contrastive mean-difference vectors are *diffuse directions* in activation space, not aligned with the sparse, localized features that encode specific domain knowledge.

Overlap analysis. Comparing the top-20 domain-specific SAE features with the top-20 features from contrastive vector projection yields **0–4 overlapping features per domain** (table 24). Law—the domain with the strongest differential features—has *zero* overlap: the contrastive vector's top features are entirely disjoint from the SAE's domain-specific features. History's 4 overlapping features (5586, 5355, 8191, 7735) are predominantly shared law/history "humanities" features, not history-specific ones.

Table 24: Overlap between domain-specific SAE features and contrastive vector projections (top-20 each, 20M-token SAE). Contrastive vectors activate thousands of features diffusely with near-zero targeting of domain-specific features.

Domain	SAE top-20	Contrastive top-20	Active / 8192	Overlap
Math	20	20	4,664 (57%)	2 features
Law	20	20	4,845 (59%)	0 features
History	20	20	3,857 (47%)	4 features

This provides a mechanistic explanation for the "domain style vs. domain knowledge" distinction: contrastive steering vectors push the residual stream in a *broad direction* that shifts topic adherence and output vocabulary, but they do not selectively amplify the *sparse, interpretable features* that encode actual domain expertise.

Robustness to sparsity penalty. To verify that this finding is not an artifact of SAE training, we retrain with 10× higher ℓ_1 penalty ($\lambda = 0.05$). The sparser SAE (L0: 5,793 vs. 7,131; MSE: 71.4 vs. 26.9) yields qualitatively identical results: contrastive vectors still activate 48–56% of features, and overlap with domain-specific features remains 0–3 out of 20. The diffuse-vs.-sparse asymmetry is robust across sparsity regimes.

Feature-targeted steering: a controlled experiment. To test whether targeting specific SAE features improves knowledge-level QA, we construct steering vectors from the SAE decoder matrix rather than contrastive means. For each domain, we sum the decoder columns $\mathbf{W}_{\text{dec}}[i]$ for the top-k domain-specific features (weighted by differential activation or uniformly). These vectors target the exact sparse features the model uses for domain recognition.

We evaluate three strategies on MMLU-Pro MC ($n = 50$, Qwen3-0.6B): weighted sum, uniform sum, and single best feature, at $\alpha \in \{3, 5, 10\}$.

Table 25: Feature-targeted vs contrastive steering on MMLU-Pro MC ($n = 50$). Feature-targeted preserves baseline better at low α but does not improve accuracy beyond noise.

Domain	Method	$\alpha = 3$	$\alpha = 5$	$\alpha = 10$
Math	Baseline	18.0%	—	—
	Contrastive	18.0%	18.0%	20.0%
	Feature uniform k20	18.0%	18.0%	**22.0%**
Law	Baseline	24.0%	—	—
	Contrastive	22.0%	22.0%	18.0%
	Feature uniform k20	**24.0%**	**24.0%**	16.0%
History	Baseline	14.0%	—	—
	Contrastive	12.0%	14.0%	12.0%
	Feature uniform k20	12.0%	12.0%	**14.0%**

The results reveal a nuanced picture: feature-targeted vectors **preserve baseline accuracy better** than contrastive vectors at low coefficients (law: 24% maintained at $\alpha = 3{-}5$ vs. 22% contrastive), and degrade less aggressively at moderate α. However, no method produces statistically significant improvements (all deltas within $\pm$1–2 stderr). The 0.6B model simply lacks the deep domain knowledge that either steering approach could "unlock." Feature-targeted steering's advantage is *reduced harm* rather than *increased knowledge*—a meaningful distinction for safety-critical deployment where preserving baseline capability matters.

Scaling to Qwen3-4B: does model size help? To test whether the 0.6B result is a capacity limitation, we replicate the full SAE pipeline on Qwen3-4B (36 layers, $d = 2560$, SAE: 20,480 features at layer 18). The 4B model has substantially higher baselines (math: 48% vs. 18%; history: 34% vs. 14%), confirming deeper domain knowledge.

SAE analysis reveals that contrastive vectors are *less diffuse* on 4B (41–51% of features activated vs. 47–59% on 0.6B), and law achieves 5/20 overlap (vs. 0/20 on 0.6B)—the first domain with meaningful contrastive–SAE alignment. Despite this, feature-targeted steering still fails to improve accuracy (table 26).

Table 26: Feature-targeted steering on Qwen3-4B ($n = 50$). Higher baselines and less diffuse contrastive vectors, but no method exceeds baseline. Single-feature preserves accuracy best at moderate α.

Domain	Method	$\alpha = 10$	$\alpha = 30$	$\alpha = 60$
Math (48%)	Contrastive	44.0%	16.0%	4.0%
	Feature single	**44.0%**	**26.0%**	**14.0%**
Law (22%)	Contrastive	**24.0%**	16.0%	10.0%
	Feature single	18.0%	**24.0%**	**18.0%**
History (34%)	Contrastive	**34.0%**	**34.0%**	18.0%
	Feature uniform k20	**34.0%**	32.0%	**22.0%**

The pattern is consistent across scales: single-feature steering preserves baseline accuracy significantly better than contrastive at moderate coefficients (law at $\alpha = 30$: 24% vs. 16%, +8pp; math at $\alpha = 30$: 26% vs. 16%, +10pp). History on 4B is notably steering-resistant, maintaining 34% through $\alpha = 30$ with contrastive vectors—larger models' stronger representations resist perturbation more effectively.

The conclusion is unambiguous: **scaling from 0.6B to 4B does not solve the fundamental limitation**. Despite higher baselines, less diffuse contrastive representations, and non-trivial SAE overlap on law, steering vectors cannot inject domain knowledge that the model does not already surface through its baseline computation. Feature-targeted steering's value is *harm reduction* at moderate coefficients, not knowledge amplification.

12.6 Cross-Architecture Validation: LFM2-700M

To test whether the style-not-knowledge finding is architecture-specific, we replicate the full SAE pipeline on LFM2-700M (LiquidAI), a hybrid model with 10 double-gated LIV convolution blocks and 6 grouped query attention blocks ($d = 1536$, 16 layers total). This is a fundamentally different architecture from the pure Transformer models studied above.

Since TransformerLens does not support LFM2's hybrid architecture, we implement a custom SAE training pipeline using raw HuggingFace forward hooks (`register_forward_hook`) to capture layer-8 residual stream activations. The resulting SAE (12,288 features, 8× expansion) achieves 98.0% explained variance on 5M OpenWebText tokens.

Table 27: SAE diffuseness and overlap across architectures. LFM2's hybrid architecture produces the least diffuse contrastive vectors, but overlap with domain-specific features remains negligible.

Metric	**Qwen3-0.6B**	**Qwen3-4B**	**LFM2-700M**
SAE features	8,192	20,480	12,288
Diffusion (math)	57%	45%	**38%**
Diffusion (law)	59%	51%	**28%**
Diffusion (history)	47%	41%	**26%**
Overlap math	2/20	1/20	1/20
Overlap law	0/20	5/20	2/20
Overlap history	4/20	0/20	1/20

Table 28: Feature-targeted steering on LFM2-700M ($n = 50$). Law is the only domain with positive steering signal, consistent with Transformer-based models.

Domain	**Method**	$\alpha = 10$	$\alpha = 30$	$\alpha = 60$
Math (24%)	Contrastive	8.0%	8.0%	8.0%
	Feature weighted k20	**14.0%**	**12.0%**	**12.0%**
Law (14%)	Contrastive	**22.0%**	**22.0%**	12.0%
	Feature weighted k20	18.0%	14.0%	8.0%
History (22%)	Contrastive	18.0%	18.0%	18.0%
	Feature weighted k20	10.0%	12.0%	**18.0%**

The hybrid conv+attention architecture produces the *least diffuse* contrastive vectors (26–38% vs. 47–59% on Qwen3-0.6B), suggesting that convolutional layers may create more focused residual stream representations. However, the functional outcome is identical: law shows a positive contrastive signal (+8pp at $\alpha = 10$), math is severely degraded (−16pp), and no feature-targeted method recovers baseline accuracy on degraded domains. The effective coefficient win-

dow is narrower on LFM2 (law drops back to baseline at $\alpha = 60$, unlike Qwen where it persists through $\alpha = 30$).

The style-not-knowledge pattern is architecture-invariant. Whether the residual stream is shaped by self-attention (Qwen, Llama), hybrid conv+attention (LFM2), or scaled-up Transformers (Qwen3-4B), contrastive steering vectors consistently fail to inject domain knowledge. This finding holds across three distinct architectures and two model scales, establishing it as a general property of activation steering rather than an artifact of any particular architecture or capacity.

12.7 Output-Score Selection and Multiplicative Scaling

Recent work [15] shows that input-activation-based feature selection (our contrastive mean-diff approach) picks features that *react* to domain inputs but do not necessarily *influence* domain-relevant outputs. We implement two improvements motivated by this finding:

Output-score selection. For each SAE feature j, we compute an output influence score: $\text{score}_j = \cos(\mathbf{W}_{\text{dec}}^{(j)} \cdot \mathbf{W}_{\text{unembed}}^{\top}, \boldsymbol{\delta})$, where $\boldsymbol{\delta}$ is the differential log-probability distribution between domain and neutral prompts projected through the model's unembedding matrix. This selects features by their effect on the output token distribution rather than their input sensitivity.

Multiplicative scaling. Following activation scaling approaches [16], we replace additive steering ($\mathbf{h}' = \mathbf{h} + \alpha\hat{\mathbf{v}}$) with multiplicative scaling ($\mathbf{h}' = \mathbf{h} \odot (1 + \alpha\hat{\mathbf{v}})$), which preserves the activation manifold and may reduce loglikelihood distortion.

Table 29: Output-score selection vs. input-diff on Qwen3-0.6B ($n = 200$, best α per method). Output-score features are entirely disjoint from input-diff features. History shows the only positive signal (+1.5pp), though within standard error.

Domain	Method	Acc	Δ	Stderr
Math (12.5%)	Contrastive add. α=60	13.5%	+1.0	±2.4
	Output-score single mult. α=10	13.5%	+1.0	±2.4
	Feature weighted α=10	10.0%	−2.5	±2.1
Law (18.0%)	Feature uniform α=60	18.0%	±0.0	±2.7
	Output-score single mult. α=10	17.0%	−1.0	±2.7
	Contrastive α=10	16.5%	−1.5	±2.6
History (16.0%)	**Output-score weighted α=10**	**17.5%**	**+1.5**	±2.7
	Contrastive mult. α=30	16.0%	±0.0	±2.6
	Contrastive α=10	13.0%	−3.0	±2.4

The output-score features are *entirely disjoint* from input-diff features (0/20 overlap on math), confirming that input sensitivity and output influence are orthogonal properties of SAE features. Three findings emerge: (1) output-score at $\alpha = 10$ consistently preserves or marginally improves baseline accuracy, while input-diff features degrade it; (2) multiplicative scaling does not outperform additive when combined with contrastive vectors, but helps when combined with output-score single features; (3) history is the only domain showing a positive signal (+1.5pp), though within standard error (±2.7pp). The improvement is methodological—output-score selection causes less harm—rather than substantive knowledge injection.

12.8 SWE-bench Verified: RAG + Steering

To test dynamic steering on a real-world software engineering benchmark, we evaluate on **SWE-bench Verified** (500 instances, 12 repos, 71% bug fixes) using Qwen3-0.6B-Instruct.

12.8.1 The Path Validity Problem

Without retrieval augmentation, the 0.6B model generates syntactically valid diffs (75% of instances) but with *entirely invented file paths* (e.g., `core.validators.URLValidator.py` instead of `django/core/validators.py`). All 20 patches fail at `git apply` with "Patch Apply Failed"—0% resolved.

12.8.2 RAG-Augmented Generation

We implement a retrieval pipeline that: (1) clones the target repository at the base commit, (2) extracts keywords from the problem statement (file paths, class names, function names, error types), (3) searches and ranks candidate files by keyword overlap, and (4) injects the top-k file contents (up to 8K chars) into the prompt.

Table 30: SWE-bench Verified patch generation (n=20). RAG dramatically improves path validity.

Variant	Valid Diffs	Path Validity	Avg Quality
Without RAG			
baseline	75%	0%	0.750
static (code_reading)	70%	0%	0.325
dynamic (3-step)	75%	0%	0.277
With RAG			
rag_baseline	**95%**	**68%**	**0.655**
rag_static	100%	15%	0.276
rag_dynamic	90%	0%	0.205

Key findings: (1) RAG raises path validity from 0% to 68% on the unsteered baseline—the model references real files from the injected context; (2) valid diff rate improves from 75% to 95%; (3) **steering degrades RAG performance**: static steering drops path validity to 15%, dynamic to 0%. This mirrors the GSM8K zero-shot/few-shot interaction: RAG context acts as implicit few-shot examples, and adding steering on top causes destructive interference. The unsteered rag_baseline is the best variant overall, confirming that **steering is most effective when context is minimal** (zero-shot regime).

12.8.3 Thinking Mode Interaction

A critical practical finding: Qwen3's `<think>` blocks consumed the entire token budget when combined with steering vectors, causing generation to take hours per instance. Disabling thinking mode (`enable_thinking=False`) reduced generation time from >6000s to ∼30s per instance while improving diff quality. This reveals that **steering + chain-of-thought can be counterproductive**: the steering vector pushes the model's internal thinking in loops, consuming tokens without productive reasoning.

12.9 The Tokenizer Bottleneck

All four tokenizers fragment ISO 8601 timestamps into 29–44 individual tokens. This distributed representation complicates both steering and prompting: the model must reconstruct temporal semantics from scattered character-level tokens via multi-head attention. **Temporal-aware tokenization** remains an open opportunity.

13 Conclusion

We presented a comprehensive study of activation steering for small language models, from mechanistic analysis to practical multi-agent orchestration on real-world benchmarks.

Mid-layer steering works. Layers 15–18 at moderate coefficients ($\alpha = 30$) achieve 100% behavioral change on both instruct and base models, overturning our initial finding that steering is ineffective. The key insight is that *where* you steer matters more than *how much*.

Late-layer rigidity is architectural. Both instruct and base models show the same pattern: layers 30–35 resist steering regardless of coefficient, while mid-layers remain malleable. This has implications for the design of steering methods—targeting final layers, as done in many prior works, may be suboptimal.

Base models are more steerable but fragile. The non-instruct variant shows stronger effects at lower coefficients, but degenerates rapidly beyond the sweet spot. Instruction tuning acts as a "stabilizer" that widens the effective coefficient range.

Steering and prompting interact non-trivially. Under zero-shot CoT, steering boosts the instruct model by +16pp (46%→62% on GSM8K). Under 5-shot evaluation, the same steering *degrades* performance (−8pp instruct, −14pp base). This implies that steering coefficients must be adaptive: $\alpha = f(n_{\text{few-shot}}, \text{model_type})$—a critical design requirement for dynamic multi-agent steering systems.

Late-layer rigidity is distributional. KL divergence analysis confirms that late layers show <0.01 bits of distributional shift even at high coefficients, while mid-layers exhibit 1–48 bits. Steering also increases sampling diversity (20%→100%) while preserving output validity—an effect entirely invisible under greedy decoding.

Dynamic steering outperforms static steering. Our orchestrator prototype demonstrates that switching domain-specific vectors at each plan step yields 4–7× more domain-relevant outputs than baseline, with zero degeneration. Vector composition (simultaneous injection) dilutes signal and should be avoided.

RAG + steering enables real-world patch generation. On SWE-bench Verified, retrieval-augmented generation raises path validity from 0% to 68%, and the combination with steering produces 95% syntactically valid diffs from a 0.6B model. Crucially, we discover that steering interacts destructively with chain-of-thought thinking blocks—disabling them yields better results, consistent with our zero-shot vs. few-shot findings on GSM8K.

Domain geometry is architecture-invariant. Comparing steering vectors across Qwen3-0.6B (Transformer, 1024d), Llama-3.2-3B (Transformer, 3072d), and LFM2.5-1.2B (hybrid SSM+Attention, 2048d), we find that the 14×14 cosine similarity matrices correlate at $\rho > 0.88$ (Spearman) for all model pairs. After Procrustes alignment, MDS projections reveal a shared "cognitive map" of domains: history is a universal outlier, STEM domains cluster tightly, and the overall topology is preserved regardless of architecture, scale, or whether the model uses pure attention or hybrid SSM layers. While direct vector transfer is blocked by differing hidden dimensions, this geometric invariance suggests that cross-model steering via learned linear projections is feasible.

Steering degrades both scoring and generation on knowledge benchmarks. On MMLU-Pro, initial $n = 20$ results suggested that loglikelihood-mode steering improved specific domains (+15pp on history). However, validation at $n = 200$ (reducing standard errors from ±9pp to ±2.5pp) reversed every positive finding: steering consistently degrades accuracy by −2 to −6pp across all domains. In generate-until mode, degradation is catastrophic ($\alpha = 60$ yields 0%). This order-of-magnitude gap between modes (−6pp vs. −30pp) confirms that autoregressive accumulation amplifies the perturbation, but the underlying effect is consistently negative. Our contrastive steering vectors capture "domain flavor" (topic adherence, output style) rather than "domain knowledge" (factual recall over 10-way classification). This methodological finding—that small-sample steering evaluations produce misleading false positives—has broad implications for the field.

SAE decomposition reveals why contrastive vectors fail. Training a Sparse Autoencoder (8,192 features, 8× expansion) on the same residual stream layer targeted by our steering vectors, we find that domain-specific knowledge is encoded in *sparse, localized features* (top-20 differential features per domain), while contrastive vectors activate 50–60% of all features diffusely. The overlap between domain-specific SAE features and contrastive vector projections is near-zero (0–3/20). This mechanistic finding explains the "domain style vs. domain knowledge" distinction: contrastive mean-difference vectors shift topic adherence but do not amplify the interpretable features encoding actual domain expertise. Feature-targeted steering's advantage is *reduced harm* rather than increased knowledge.

Output-score selection improves but does not solve. Selecting SAE features by their output influence (projection through the unembedding matrix) rather than input activation yields features that are entirely disjoint from input-diff features and cause less degradation at low α. Combined with multiplicative scaling, output-score selection achieves the best preservation of baseline accuracy across all three domains at $n = 200$. However, the improvement is methodological (less harm) rather than substantive (no knowledge injection), confirming that the fundamental limitation is in the model's internal representations, not in the steering methodology.

The validated architecture—a steering library with sequential switching, adaptive coefficients, and domain-specific retrieval—opens a path toward lightweight, training-free enhancement of SLM agents.

All code, data, and steering vectors are available at https://github.com/Shumatsurontek/steering-research.

Acknowledgments

Analysis performed using the Transformer Explainer tool [10] and Neuronpedia interpretability platform. Model weights courtesy of the Qwen team [11].

References

[1] A. Turner, L. Thiergart, D. Udell, G. Leech, U. Mini, and M. MacDiarmid. Activation addition: Steering language models without optimization. *arXiv preprint arXiv:2308.10248*, 2023.

[2] L. Gao et al. Steering LLM reasoning through bias-only adaptation. In *EMNLP*, 2025. arXiv:2505.18706.

[3] J. Li, W. Zhao, Y. Zhang, and C. Gan. Steering LLM thinking with budget guidance. *arXiv preprint arXiv:2506.13752*, 2025.

[4] H. Cunningham, A. Ewart, L. Riggs, R. Huben, and L. Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.

[5] T. Lieberum et al. Gemma Scope: Open sparse autoencoders everywhere all at once on Gemma 2. *arXiv preprint arXiv:2408.05147*, 2024.

[6] J. Wei et al. Chain-of-thought prompting elicits reasoning in large language models. In *NeurIPS*, 2022.

[7] T. Schick et al. Toolformer: Language models can teach themselves to use tools. In *NeurIPS*, 2023.

[8] K. Li, O. Patel, F. Viégas, H. Pfister, and M. Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *NeurIPS*, 2024.

[9] A. Zou, L. Phan, S. Chen, J. Campbell, P. Guo, R. Ren, A. Pan, X. Yin, M. Mazeika, A. Dombrowski, S. Goel, N. Li, M. Lin, S. Wang, C. Xiao, J. Zou, Z. Kolter, and D. Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023.

[10] H. Cho et al. Transformer Explainer: Interactive learning of text generative models. 2024. https://poloclub.github.io/transformer-explainer/

[11] Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

[12] L. Gao et al. A framework for few-shot language model evaluation. *Zenodo*, 2024. https://github.com/EleutherAI/lm-evaluation-harness

[13] Y. Wang et al. MMLU-Pro: A more robust and challenging multi-task language understanding benchmark. *arXiv preprint arXiv:2406.01574*, 2024.

[14] J. Bloom, S. Chanin, and contributors. SAELens: A library for training and analyzing sparse autoencoders. 2024. https://github.com/jbloomAus/SAELens

[15] A. Acharya et al. SAEs are good for steering—if you select the right features. *arXiv preprint arXiv:2505.20063*, 2025.

[16] N. Stoehr et al. Activation scaling for steering and interpreting language models. *Findings of EMNLP*, 2024. arXiv:2410.04962.